

STAT452 Final Project: Wine Quality and Experimental Design

Trần Lê Anh Tuấn (24125107) Nguyễn Bảo Minh Triết (24125047)
Nguyễn Đình Thiên Lộc (24125093) Nguyễn Hồng Tấn Tài (24125078)
Lê Minh Thuận (24125105)

Academic Year 2025–2026

Contents

1	Part 1: Wine Quality (Red)	2
2	Exploratory Data Analysis	2
2.1	Dataset, variables, and prediction question	2
2.2	Response and predictor distributions	2
2.3	Pairwise relationships and correlation structure	4
3	Data Cleaning	5
3.1	Missingness, duplicates, and the locked split	5
3.2	Outliers and transformations	5
3.3	Leakage control	7
4	Feature Selection	7
4.1	Screening without deleting correlated information	7
4.2	Mean and OLS baselines	8
4.3	Embedded Lasso selection	8
5	Modelling with Regularization	8
5.1	Cross-validation and shrinkage	9
5.2	Pre-holdout decision	10
6	Evaluation	10
6.1	Fair model comparison	10
6.2	Diagnostics and bias–variance trade-off	12
6.3	Limitations	12
6.4	Conclusion	12
7	Part 2: Experimental-design extension	13
8	Contributions	13
A	Detailed data documentation	16
B	Cleaning and outlier audit	16
C	Supplementary selection and model tables	16
D	Reproducibility and collaboration	16
	References	18

Abstract

This report currently completes Part 1 of the Group 8 final project. We predict red Vinho Verde quality from 11 physicochemical measurements with a locked, leakage-controlled comparison of OLS, Ridge, Lasso, and Elastic Net. Part 2, the separate experimental-design extension, remains explicitly pending.

1 Part 1: Wine Quality (Red)

2 Exploratory Data Analysis

2.1 Dataset, variables, and prediction question

The UCI Wine Quality data describe red Portuguese Vinho Verde samples using 11 continuous physicochemical measurements and a sensory score supplied as the median of at least three expert ratings (Cortez et al. 2009a, 2009b). The response is **quality**; the predictors cover acidity, sugar, chlorides, sulfur dioxide, density, pH, sulphates, and alcohol. All predictors are numeric, and the detailed data dictionary appears in Appendix A.

We model quality numerically because regression preserves its ordering and avoids an arbitrary good/poor threshold. This assumes that a one-point change has roughly comparable meaning across the observed 3–8 range. Gaussian regression can produce fractional expected scores, so the interpretation is prediction of expected expert score rather than exact ordinal classification.

After the integrity checks in Section 2, the unique records were divided by a fixed quality-stratified 80/20 split. Every plot below uses only the 1088 training observations; the 271 holdout outcomes were not inspected.

2.2 Response and predictor distributions




Figure 1: Training-set distribution of the sensory response. Scores 5 and 6 dominate, while the extremes contain few observations.

The response is concentrated at 5 and 6. A mean-like prediction is therefore a strong baseline, and error estimates for rare scores 3, 4, and 8 will be less precise. This imbalance also foreshadows regression toward the center.




Figure 2: Training-set predictor distributions before transformation. Several chemistry measures are strongly right-skewed.

Residual sugar, chlorides, sulfur-dioxide measures, and sulphates have long right tails. These distributions motivate the transformations documented in Section 2; EDA itself is shown on the original measurement scale.

2.3 Pairwise relationships and correlation structure




Figure 3: Training-set Pearson correlations. Blocks among acidity, density, pH, and sulfur-dioxide measures motivate coefficient stabilization.

The strongest marginal response signals are alcohol ($r = 0.47$); volatile acidity ($r = -0.37$); sulphates ($r = 0.23$); citric acid ($r = 0.22$). The heatmap also shows correlated predictor blocks, especially among acidity, density, pH, and the two sulfur-dioxide measurements. That structure makes OLS coefficients potentially unstable and directly motivates Ridge, Lasso, and Elastic Net.




Figure 4: The four strongest marginal predictor–quality relationships in training. Jitter reveals the discrete response; lines are descriptive OLS trends.

Alcohol is positively associated with quality and volatile acidity negatively associated, but every score overlaps substantially on every predictor. No single physicochemical measurement is sufficient. These are predictive associations, not causal effects, because grape variety, brand, vintage, price, and production conditions are absent.

3 Data Cleaning

3.1 Missingness, duplicates, and the locked split

The raw archive contains 1599 rows, 11 numeric predictors, no missing cells, and no non-finite values. It also contains 240 exact duplicate records. Because the archive provides no sample identifier, an exact repeat could be either a copied row or a genuinely repeated batch. We conservatively keep the first copy and remove the rest **before** splitting, preventing identical profiles and labels from appearing in both training and holdout data. The limitation is that this choice may underweight genuinely repeated batches.

The fixed seed 4520803 produces a quality-stratified 80/20 partition after deduplication. Seed 4520804 then assigns five stratified folds inside training. The shared fold IDs are reused by every model.

3.2 Outliers and transformations

Tukey’s 1.5-IQR rule flags 277 training rows (25.5%). Deleting all of them would discard a large and scientifically plausible part of the target population; the dataset metadata also warns that rare wines are expected. We therefore

Table 1: Data-integrity audit and decisions.

Check	Result	Decision
Raw observations	1599	Validated against UCI metadata
Numeric predictors	11	All retained
Missing cells	0	None to impute
Non-finite numeric cells	0	None
Exact duplicate rows removed	240	Remove before splitting to prevent duplicate leakage
Unique observations	1359	Analysis population
Observed quality range	3–8	Treat as numeric regression response

Table 2: Locked data partitions after deduplication.

Partition	Observations	Percent_of_unique
Raw archive	1599	NA
Unique before split	1359	100.0
Training	1088	80.1
Holdout	271	19.9

retain every row, winsorize each predictor at its training 1st and 99th percentiles to limit leverage, and never clip the response.

Outlier flags are common and distributed across variables




Figure 5: Training observations flagged by the univariate Tukey-IQR rule. Flags are diagnosed and retained rather than deleted wholesale.

The training skewness rule selects residual sugar, chlorides, free sulfur dioxide, total sulfur dioxide, sulphates for `log1p` transformation. Median imputation is defined as a reproducibility safeguard but is a no-op for this dataset. After winsorization and optional `log1p`, all predictors are centered and scaled so penalty magnitudes are comparable.

3.3 Leakage control

All medians, winsorization limits, transformation choices, means, and standard deviations are learned without holdout outcomes. During cross-validation, those quantities are refit inside each analysis fold and applied to its validation fold. The final recipe is fit once to all training predictors and then applied unchanged to the predictor-only holdout file. This design prevents preprocessing information from leaking backward into feature selection or tuning.

4 Feature Selection

4.1 Screening without deleting correlated information

We use three training-only views of relevance: marginal correlations, variance inflation factors (VIFs), and an AIC stepwise sensitivity analysis. The largest VIF is 7.90 for fixed acidity. This is material multicollinearity but not evidence that the feature is useless: deleting one correlated chemistry measure can make the retained coefficient depend arbitrarily on that choice. Therefore all predictors enter Ridge, while sparse selection is delegated to Lasso’s embedded penalty.




Figure 6: VIFs and full-OLS coefficients after the training-derived transformations. Coefficients are comparable because predictors are standardized.

The AIC sensitivity model retains 7 of 11 predictors: alcohol, volatile acidity, sulphates, chlorides, ph, total sulfur dioxide, free sulfur dioxide. We do not promote this data-dependent subset directly to the holdout comparison; it is a diagnostic against which to compare Lasso.

Table 3: Mean and OLS baseline performance before holdout evaluation.

Model	Training_RMSE	CV_RMSE	CV_MAE	CV_R2	Predictors
Mean	0.824	0.824	0.695	0.000	0
OLS	0.663	0.673	0.516	0.334	11

Table 4: Embedded Lasso feature selection using the one-standard-error penalty.

Feature	Coefficient	Selected
alcohol	0.2996	Yes
volatile_acidity	-0.1664	Yes
sulphates	0.1253	Yes
chlorides	-0.0360	Yes
ph	-0.0177	Yes
total_sulfur_dioxide	-0.0136	Yes
fixed_acidity	0.0038	Yes
citric_acid	0.0000	No
residual_sugar	0.0000	No
free_sulfur_dioxide	0.0000	No
density	0.0000	No

4.2 Mean and OLS baselines

The mean-only model establishes the no-chemistry benchmark. OLS lowers cross-validated error, but correlated predictors allow coefficient variance. Regularization deliberately accepts some training bias to reduce that variance.

4.3 Embedded Lasso selection

At the one-standard-error penalty, Lasso retains 7 predictors: alcohol, volatile acidity, sulphates, chlorides, ph, total sulfur dioxide, fixed acidity. This is the required principled feature-selection result. The one-SE rule is used for interpretation because it favors a stable sparse description; the minimum-CV penalty remains the predeclared predictive fit.

5 Modelling with Regularization

For standardized transformed predictors, all regularized candidates minimize

$$\frac{1}{2n} \|\mathbf{y} - \beta_0 - \mathbf{X}\beta\|_2^2 + \lambda \left\{ \alpha \|\beta\|_1 + \frac{1-\alpha}{2} \|\beta\|_2^2 \right\}.$$

Ridge uses $\alpha = 0$ (Hoerl and Kennard 1970), Lasso $\alpha = 1$ (Tibshirani 1996), and Elastic Net searches $\alpha \in \{0.25, 0.50, 0.75\}$ (Zou and Hastie 2005). For each alpha, 120 decreasing λ values from 100 to 10^{-4} use identical folds. Coordinate-descent fits are produced with `glmnet` (Friedman et al. 2010).

Table 5: Regularized-model tuning and effective size on the shared training folds.

Model	Alpha	Lambda_min	Lambda_1se	CV_RMSE_min	Nonzero_min	Nonzero_1se
Ridge	0.00	0.0528	0.2683	0.6713	11	11
Lasso	1.00	0.0052	0.0373	0.6712	9	7
Elastic Net	0.75	0.0065	0.0470	0.6712	9	7

5.1 Cross-validation and shrinkage




Figure 7: Five-fold training CV curves with fold-specific preprocessing. Dashed lines minimize CV error; dotted lines apply the one-standard-error rule.

Elastic Net selected $\alpha = 0.75$. The minimum-error and one-SE penalties, CV errors, and effective model sizes are reported below. Ridge keeps all coefficients but continuously shrinks them; Lasso and Elastic Net can set coefficients exactly to zero.

Table 6: Fair holdout comparison; the model lock was based only on five-fold CV RMSE.

Model	CV_RMSE	Holdout_RMSE	RMSE_CI_Lower	RMSE_CI_Upper	Holdout_MAE	Holdout_R2	Locked_by_CV
OLS	0.6726	0.6131	0.5565	0.6677	0.4803	0.4406	No
Ridge	0.6713	0.6135	0.5571	0.6717	0.4824	0.4397	No
Elastic Net	0.6712	0.6139	0.5582	0.6703	0.4822	0.4391	No
Lasso	0.6712	0.6139	0.5582	0.6703	0.4822	0.4390	Yes
Mean	0.8242	0.8198	0.7461	0.9018	0.6923	-0.0001	No




Figure 8: Coefficient paths as regularization changes. The vertical line in each panel is the minimum-CV penalty used for prediction.

As λ grows, paths move toward zero. Lasso reaches exact zeros earliest; Ridge shares weight among correlated measurements; Elastic Net interpolates. This is the expected bias–variance mechanism rather than merely a software choice.

5.2 Pre-holdout decision

Before reconstructing any holdout outcome, we locked **Lasso**, whose five-fold CV RMSE was 0.6712. This written lock prevents choosing a preferred model after seeing test performance. All predeclared models are nevertheless evaluated to make the comparison transparent.

6 Evaluation

6.1 Fair model comparison

On the same 271 rows, the pre-locked Lasso obtained RMSE 0.614, MAE 0.482, and R^2 0.439. The observed lowest holdout RMSE belongs to **OLS** (0.613). We retain the CV lock as the confirmatory choice even if this ranking differs: selecting by the holdout would make its error optimistic. The bootstrap intervals are wide enough that small ranking differences should not be overinterpreted.




Figure 9: Held-out predictions from the four fitted regression models. The diagonal denotes perfect prediction.

Predictions shrink toward the common scores 5–6, which lowers average error but systematically misses rare extremes. This pattern follows from response imbalance and squared-error fitting, not simply from regularization.

6.2 Diagnostics and bias–variance trade-off




Figure 10: Residual checks for the pre-locked model on the untouched holdout set. These diagnostics describe generalization and were not used to retune the model.

OLS has the least shrinkage and therefore the lowest regularization bias but the greatest sensitivity to correlated predictors. Ridge reduces coefficient variance without selection; Lasso trades additional bias for sparsity; Elastic Net stabilizes groups of correlated features while retaining some selection. The relevant empirical comparison is the gap among training, CV, and holdout RMSE in Table 6. The regularized candidates have similar held-out accuracy, so stability and interpretability matter more than a small decimal-place difference.

6.3 Limitations

1. Quality is ordinal and discrete; Gaussian regression treats adjacent scores as equally spaced and permits fractional predictions.
2. Rare scores 3, 4, and 8 yield imprecise tail performance, while the bootstrap quantifies only sampling variation in this holdout.
3. Exact deduplication prevents leakage but may remove genuinely repeated batches because the archive has no sample identifier.
4. Physicochemical measurements omit brand, vintage, grape variety, price, and production process; predictive associations are not causal effects.
5. Winsorization limits leverage and log transforms reduce skew, but alternative robust losses or ordinal models could change tail behavior.

6.4 Conclusion

The full workflow meets the Part 1 objective: documented EDA and cleaning, principled selection, a baseline plus three regularized models, shared-fold CV, coefficient-path interpretation, and a single locked holdout evaluation. Lasso is the confirmatory recommendation for expected-score prediction under this design. Alcohol and volatile acidity provide the strongest marginal signals, but correlated chemistry measurements and irreducible sensory variation limit accuracy. A future extension should compare ordinal regression or robust nonlinear models under the same locked resampling protocol.

7 Part 2: Experimental-design extension

TODO: Select and justify the separate dataset; state hypotheses; complete the two-factor ANOVA or 2^k factorial analysis; check assumptions; report interactions and post-hoc comparisons where justified; and interpret practical significance and limitations.

8 Contributions

TODO: Replace the workstream ownership plan with a statement of the work each member actually completed and reviewed.

Table 7: Wine Quality (Red) data dictionary.

Variable	Role	Type	Description
fixed_acidity	Predictor	Numeric	Fixed acidity (g/dm3)
volatile_acidity	Predictor	Numeric	Volatile acidity (g/dm3)
citric_acid	Predictor	Numeric	Citric acid (g/dm3)
residual_sugar	Predictor	Numeric	Residual sugar (g/dm3)
chlorides	Predictor	Numeric	Chlorides (g/dm3)
free_sulfur_dioxide	Predictor	Numeric	Free sulfur dioxide (mg/dm3)
total_sulfur_dioxide	Predictor	Numeric	Total sulfur dioxide (mg/dm3)
density	Predictor	Numeric	Density (g/cm3)
ph	Predictor	Numeric	pH
sulphates	Predictor	Numeric	Potassium sulphate (g/dm3)
alcohol	Predictor	Numeric	Alcohol (% by volume)
quality	Response	Integer score	Median sensory score from wine experts (0–10 documented; 3–8 observed)

Table 8: Training-set descriptive statistics.

Variable	N	Mean	SD	Min	Q1	Median	Q3	Max	Skewness
fixed_acidity	1088	8.259	1.713	4.600	7.100	7.850	9.100	15.900	0.955
volatile_acidity	1088	0.529	0.181	0.120	0.390	0.520	0.640	1.580	0.755
citric_acid	1088	0.271	0.194	0.000	0.090	0.260	0.420	1.000	0.321
residual_sugar	1088	2.558	1.445	1.200	1.900	2.200	2.600	15.500	4.453
chlorides	1088	0.089	0.052	0.012	0.070	0.079	0.091	0.611	5.448
free_sulfur_dioxide	1088	16.156	10.638	1.000	8.000	14.000	22.000	72.000	1.265
total_sulfur_dioxide	1088	47.348	33.716	6.000	22.000	38.000	64.250	289.000	1.569
density	1088	0.997	0.002	0.990	0.996	0.997	0.998	1.004	0.052
ph	1088	3.312	0.155	2.740	3.210	3.310	3.400	4.010	0.196
sulphates	1088	0.661	0.176	0.370	0.550	0.620	0.730	2.000	2.534
alcohol	1088	10.428	1.088	8.400	9.500	10.100	11.100	14.900	0.876
quality	1088	5.625	0.825	3.000	5.000	6.000	6.000	8.000	0.198

Table 9: Training-set Tukey-IQR outlier audit; rows are retained.

Variable	IQR_Flags	Percent	Lower_Fence	Upper_Fence
fixed_acidity	34	3.125	4.100	12.100
volatile_acidity	16	1.471	0.015	1.015
citric_acid	1	0.092	-0.405	0.915
residual_sugar	109	10.018	0.850	3.650
chlorides	75	6.893	0.039	0.122
free_sulfur_dioxide	22	2.022	-13.000	43.000
total_sulfur_dioxide	31	2.849	-41.375	127.625
density	29	2.665	0.992	1.001
ph	24	2.206	2.925	3.685
sulphates	46	4.228	0.280	1.000
alcohol	10	0.919	7.100	13.500

Table 10: Training-only transformation and winsorization plan.

Variable	Raw_Skewness	IQR_Flags	Winsor_Lower	Winsor_Upper	Transformation
fixed_acidity	0.9547	34	5.1000	12.9000	standardize
volatile_acidity	0.7550	16	0.2000	1.0269	standardize
citric_acid	0.3214	1	0.0000	0.7013	standardize
residual_sugar	4.4527	109	1.4000	8.6260	log1p + standardize
chlorides	5.4480	75	0.0420	0.3888	log1p + standardize
free_sulfur_dioxide	1.2646	22	3.0000	51.0000	log1p + standardize
total_sulfur_dioxide	1.5691	31	8.0000	147.0000	log1p + standardize
density	0.0522	29	0.9919	1.0014	standardize
ph	0.1959	24	2.9387	3.7100	standardize
sulphates	2.5345	46	0.4300	1.3126	log1p + standardize
alcohol	0.8758	10	9.0000	13.4130	standardize

A Detailed data documentation

B Cleaning and outlier audit




Figure 11: Training observations flagged by the univariate Tukey-IQR rule. Flags are diagnosed rather than deleted wholesale.

C Supplementary selection and model tables

D Reproducibility and collaboration

The canonical analysis is rebuilt from RStudio with `source("analysis/00_run_all.R")`, or from a terminal with `Rscript --vanilla analysis/00_run_all.R`. The five workstreams execute in dependency order, every random operation uses a recorded seed, and `output/artifact_manifest.csv` stores file sizes and MD5 checksums. Row-level holdout outcomes are not serialized. Software versions for the verified run are listed below; full `sessionInfo()` is in `output/logs/session_info.txt`. The analysis and dynamic report use R, `knitr`, and `rmarkdown` (R Core Team 2026; Xie 2025; Allaire et al. 2026).

The workstream allocation and review interfaces are documented in the repository's collaboration work-split file. These organizational labels must be checked against the members' actual contributions before final submission.

Table 11: Training-only feature screening and OLS coefficients on the standardized scale.

Feature	Correlation_with_quality	VIF	OLS_standardized_coefficient	Stepwise_AIC_retained
alcohol	0.475	3.511		0.312 Yes
volatile_acidity	-0.369	1.829		-0.189 Yes
sulphates	0.280	1.508		0.179 Yes
citric_acid	0.220	3.225		-0.061 No
density	-0.177	6.993		-0.017 No
total_sulfur_dioxide	-0.157	3.335		-0.089 Yes
chlorides	-0.133	1.525		-0.088 Yes
fixed_acidity	0.116	7.903		0.042 No
ph	-0.050	3.409		-0.076 Yes
free_sulfur_dioxide	-0.047	3.011		0.058 Yes
residual_sugar	0.035	1.834		0.013 No

Table 12: Coefficients at minimum-CV penalties on the standardized transformed scale.

Feature	OLS	Ridge	Lasso	Elastic_Net
fixed_acidity	0.0424	0.0575	0.0159	0.0165
volatile_acidity	-0.1894	-0.1725	-0.1799	-0.1799
citric_acid	-0.0607	-0.0313	-0.0330	-0.0336
residual_sugar	0.0126	0.0213	0.0000	0.0000
chlorides	-0.0876	-0.0848	-0.0842	-0.0844
free_sulfur_dioxide	0.0577	0.0440	0.0357	0.0367
total_sulfur_dioxide	-0.0895	-0.0784	-0.0689	-0.0699
density	-0.0167	-0.0489	0.0000	0.0000
ph	-0.0756	-0.0509	-0.0707	-0.0708
sulphates	0.1792	0.1727	0.1691	0.1692
alcohol	0.3125	0.2762	0.3179	0.3175

Table 13: Elastic Net mixing-parameter search on shared folds.

Alpha	CV_RMSE_min	Lambda_min
0.25	0.6713	0.0131
0.50	0.6712	0.0093
0.75	0.6712	0.0065

Table 14: Holdout error by observed score for the pre-locked Lasso model.

Quality	N	Mean_Error	MAE	RMSE
3	2	-1.8395	1.8395	1.8439
4	11	-1.1948	1.1948	1.2366
5	115	-0.3244	0.3470	0.4526
6	107	0.1762	0.4082	0.4860
7	33	0.7751	0.7751	0.8213
8	3	1.5685	1.5685	1.5744

Table 15: Software versions used for the verified run.

Package	Version
R	4.5.2
glmnet	5.0
knitr	1.51
rmarkdown	2.31

References

- Allaire, JJ, Yihui Xie, Christophe Dervieux, et al. 2026. *Rmarkdown: Dynamic Documents for r*. <https://github.com/rstudio/rmarkdown>.
- Cortez, Paulo, António Cerdeira, Fernando Almeida, Telmo Matos, and José Reis. 2009a. “Modeling Wine Preferences by Data Mining from Physicochemical Properties.” *Decision Support Systems* 47 (4): 547–53. <https://doi.org/10.1016/j.dss.2009.05.016>.
- Cortez, Paulo, António Cerdeira, Fernando Almeida, Telmo Matos, and José Reis. 2009b. *Wine Quality*. UCI Machine Learning Repository. <https://archive.ics.uci.edu/dataset/186/wine+quality>.
- Friedman, Jerome, Trevor Hastie, and Robert Tibshirani. 2010. “Regularization Paths for Generalized Linear Models via Coordinate Descent.” *Journal of Statistical Software* 33 (1): 1–22. <https://doi.org/10.18637/jss.v033.i01>.
- Hoerl, Arthur E., and Robert W. Kennard. 1970. “Ridge Regression: Biased Estimation for Nonorthogonal Problems.” *Technometrics* 12 (1): 55–67. <https://doi.org/10.1080/00401706.1970.10488634>.
- R Core Team. 2026. *R: A Language and Environment for Statistical Computing*. R Foundation for Statistical Computing. <https://doi.org/10.32614/R.manuals>.
- Tibshirani, Robert. 1996. “Regression Shrinkage and Selection via the Lasso.” *Journal of the Royal Statistical Society: Series B* 58 (1): 267–88. <https://doi.org/10.1111/j.2517-6161.1996.tb02080.x>.
- Xie, Yihui. 2025. *Knitr: A General-Purpose Package for Dynamic Report Generation in r*. <https://yihui.org/knitr/>.
- Zou, Hui, and Trevor Hastie. 2005. “Regularization and Variable Selection via the Elastic Net.” *Journal of the Royal Statistical Society: Series B* 67 (2): 301–20. <https://doi.org/10.1111/j.1467-9868.2005.00503.x>.